


local-inference-setup

Running **Qwen3.8-Flash-Next** (176B-total / ~6B-active MoE) on a **Mac mini M4 Pro, 64 GB**: the exact setup, measured speeds, and a month of notes on what helped. Plus the agent instructions and small tools I use with it.



Current setup (October 2026)

Part	What	From
Model	Swift 1.5 GSQ-RCO IQ3_XXS	ukisai/Swift-1.5-Qwen3.8-Flash-Next-GSQ-RCO-GGUF
MTP head	MTP/mtp-Qwen3.8-Flash-Next-shared-Q8_0.gguf	unsloth/Qwen3.8-Flash-Next-GGUF
Server	llama.cpp fork b11443-mix-d65395f	unslothai/llama.cpp
Script	262k context, MTP 6 / p-min 0.8, optional vision	setup/run-swift-mtp.sh

Decode on real agent work at 64–128k context: **~25 t/s median**. Memory: 56.7 GiB wired at idle of a 60 GiB GPU limit; swap stays flat.

What’s here

Folder	Contents
posts/	2026-10: one llama.cpp upgrade beat a month of tuning (latest; Reddit body) · 2026-09: the original setup (superseded)
history/	The whole month: who supplies which part (Qwen, AtomicChat, ISTA-DASLab, UkisAI, Unsloth, ggml-org), the timeline, decode by era, quality vs BF16, tried and rejected
setup/	The run script, with download commands in its header
agents/	My global AGENTS.md for coding agents (Claude Code, pi) and the vendored PONYTAIL rules
tools/	md-table-fit: wrap over-wide Markdown tables, --join to unwrap · md2pdf: Markdown to PDF offline, Mermaid drawn

SQL tools (SQL Server 2016 query checker, T-SQL keyword-river formatter) live in [leastsurprise/sql-tools](#).

Speeds were measured on this machine; quality figures (KLD, benchmarks) are the publishers’ own.